

(1)

Kubernetes Networking Demystified: A Brief Guide



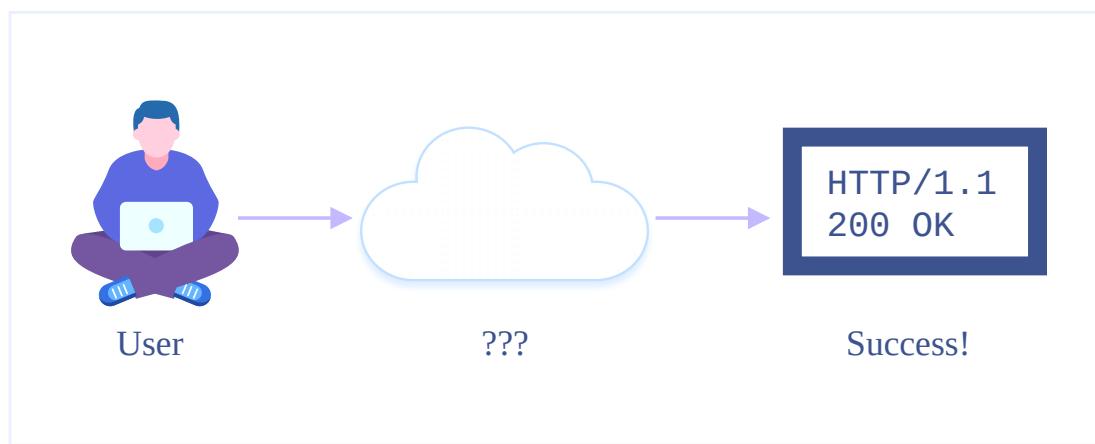
[Karen Bruner \(/authors/kbruner/\)](#) Jan 09, 2020

Kubernetes cluster networking can be more than a bit confusing, even for engineers with hands-on experience working with virtual networks and request routing. In this post, we will present an introduction into the complexities of Kubernetes networking by following the journey of an HTTP request to a service running on a basic Kubernetes cluster.

We will use a standard Google Kubernetes Engine (GKE) cluster with two Linux nodes for our examples, with notes on where the details might differ on other platforms.

A Request's Journey

Take the typical example of a person browsing the web. They click on a link, something happens, and a page loads.

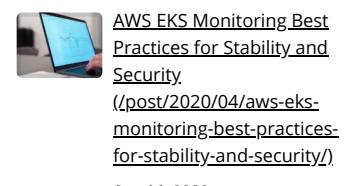


FEATURED POSTS



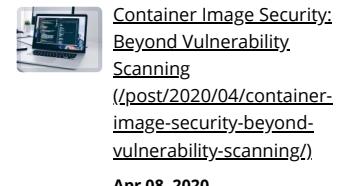
[Enhancing Kubernetes Security with Open Policy Agent \(OPA\) - Part 1](#)
(/post/2020/04/enhancing-kubernetes-security-with-open-policy-agent-opa-part-1/)

Apr 29, 2020



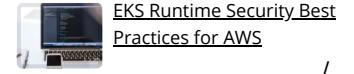
[AWS EKS Monitoring Best Practices for Stability and Security](#)
(/post/2020/04/aws-eks-monitoring-best-practices-for-stability-and-security/)

Apr 14, 2020



[Container Image Security: Beyond Vulnerability Scanning](#)
(/post/2020/04/container-image-security-beyond-vulnerability-scanning/)

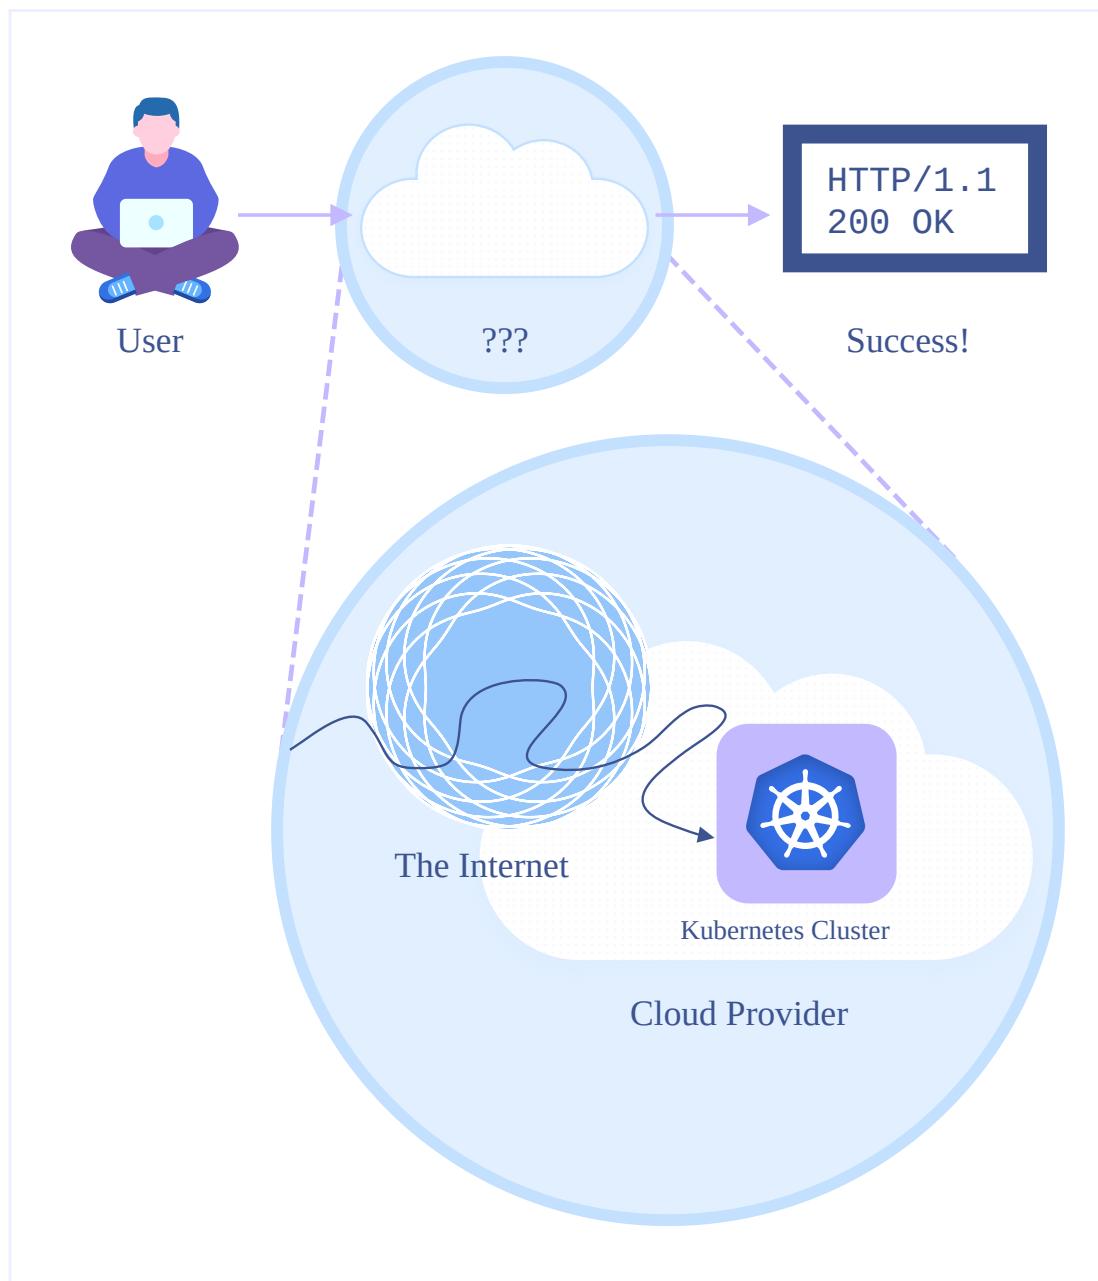
Apr 08, 2020



[EKS Runtime Security Best Practices for AWS](#)

ext diagram, the request gets sent through the Internet to a very large cloud provider, then to a

(/) kubernetes cluster hosted in the cloud provider's infrastructure.



Guide to Kubernetes Network Policies for Maximum Security

Guide to Kubernetes Network Policies

Download our step-by-step guide for setting up Kubernetes network policies

[DOWNLOAD NOW](#)

practices-for-aws-workloads/)

Mar 05, 2020

[EKS vs GKE vs AKS - April 2020 Updates](#)
(/post/2020/03/eks-vs-gke-vs-aks-april-2020-updates/)

Mar 31, 2020

[EKS Networking Best Practices for Security and Operation](#)
(/post/2020/03/eks-networking-best-practices/)

Mar 30, 2020

[Securing EKS Cluster Add-ons: Dashboard, Fargate, EC2 Components, and More](#)
(/post/2020/03/securing-eks-cluster-add-ons-dashboard-fargate-ec2-and-more/)

Mar 24, 2020

[What's New in Kubernetes 1.18? New Features and Updates](#)
(/post/2020/03/what-is-new-in-kubernetes-1.18/)

Mar 23, 2020

[Guide to Designing EKS Clusters for Better Security](#)
(/post/2020/03/guide-to-eks-cluster-design-for-better-security/)

Mar 17, 2020

[Azure Kubernetes \(AKS\) Security Best Practices Part 4 of 4: Cluster Maintenance](#)
(/post/2020/03/azure-kubernetes-aks-security-best-practices-part-4-of-4/)

Mar 09, 2020

[Azure Kubernetes \(AKS\) Security Best Practices Part 3 of 4: Runtime Security](#)
(/post/2020/02/azure-kubernetes-aks-security-best-practices-part-3-of-4/)

Feb 24, 2020

[Azure Kubernetes \(AKS\) Security Best Practices Part 2 of 4: Networking](#)
(/post/2020/02/azure-kubernetes-aks-security-best-practices-part-2-of-4/)

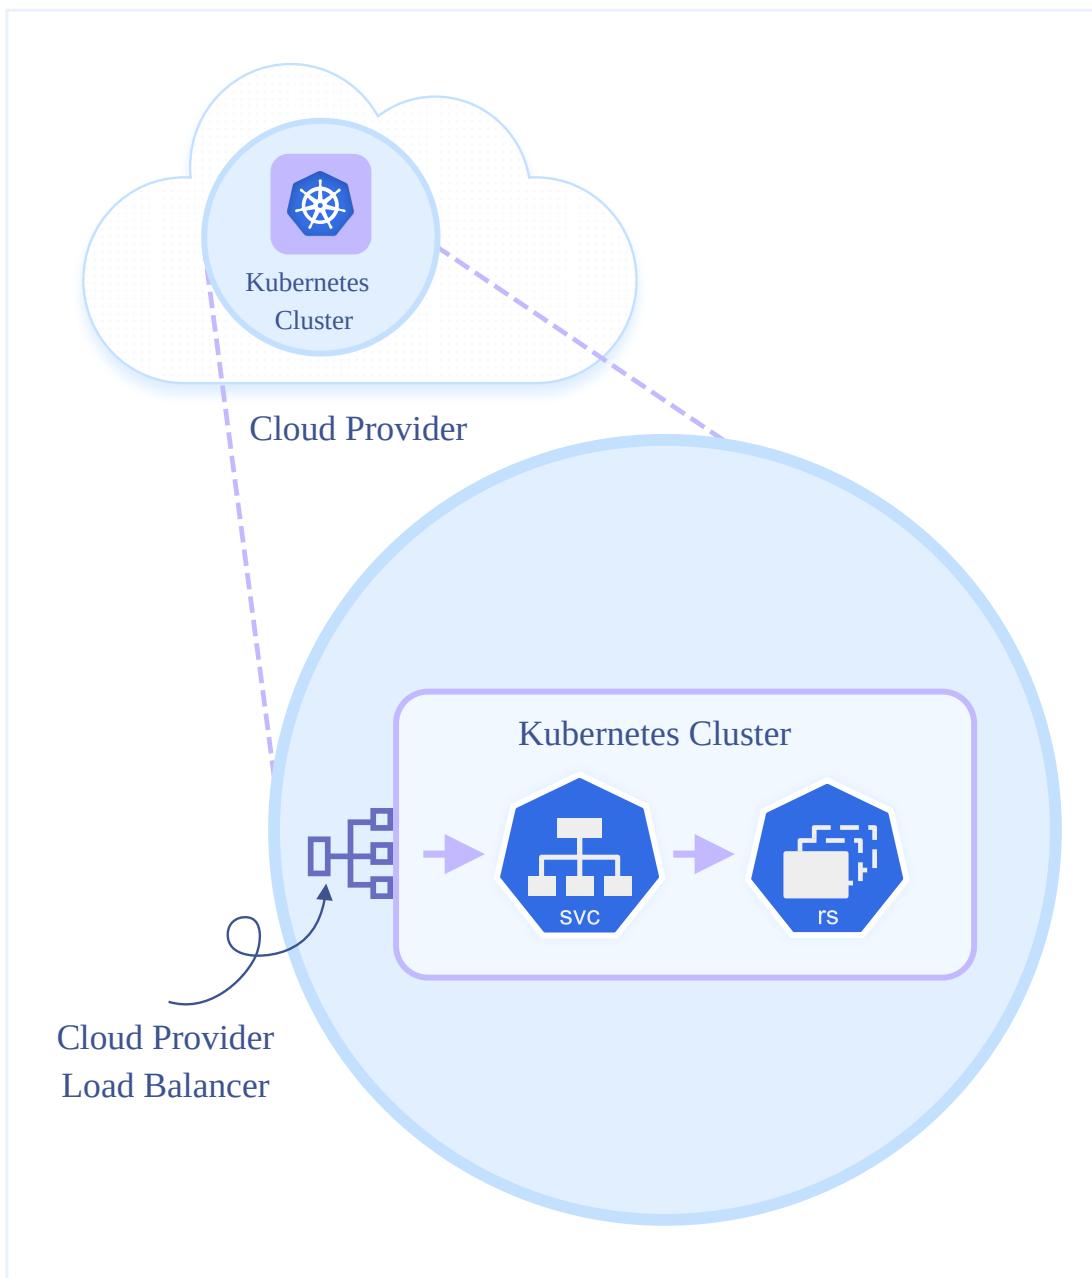
Feb 11, 2020

[EKS vs GKE vs AKS - Evaluating Kubernetes in the Cloud](#)

Feb 10, 2020

Zoom in closer to the Kubernetes cluster, we see a cloud provider load balancer feeding to a

(1) Kubernetes Service resource, which then routes requests to pods in a Kubernetes ReplicaSet.



We can deploy the following YAML to create the Kubernetes Service (svc) and ReplicaSet (rs):

-  [Azure Kubernetes \(AKS\) Security Best Practices Part 1 of 4: Designing Secure Clusters and Container Images](#)
`(/post/2020/01/azure-kubernetes-aks-security-best-practices-part-1-of-4/)`
Jan 27, 2020
-  [Guide to Kubernetes Egress Network Policies](#)
`(/post/2020/01/kubernetes-egress-network-policies/)`
Jan 15, 2020
-  [Top 5 Kubernetes Vulnerabilities of 2019 - the Year in Review](#)
`(/post/2020/01/top-5-kubernetes-vulnerabilities-of-2019-the-year-in-review/)`
Jan 02, 2020
-  [What's New in Kubernetes 1.17: A Deeper Look at New Features](#)
`(/post/2019/12/whats-new-in-kubernetes-1.17-a-deeper-look-at-new-features/)`
Dec 09, 2019
-  [How to Make Istio Work with Your Apps](#)
`(/post/2019/11/how-to-make-istio-work-with-your-apps/)`
Nov 26, 2019
-  [Protecting Kubernetes API Against CVE-2019-11253 \(Billion Laughs Attack\) and Other Vulnerabilities](#)
`(/post/2019/09/protecting-kubernetes-api-against-cve-2019-11253-billion-laughs-attack/)`
Sep 30, 2019
-  [12 Kubernetes configuration best practices](#)
`(/post/2019/09/12-kubernetes-configuration-best-practices/)`
Sep 26, 2019
-  [Docker Container Security 101: Risks and 33 Best Practices](#)
`(/post/2019/09/docker-security-101/)`
Sep 13, 2019
-  [Amazon EKS Security Best Practices](#)
`(/post/2019/09/amazon-eks-security-best-practices/)`

Sep 13, 2019

```

d: ReplicaSet
()
  metadata:
    name: hello-world
    labels:
      app: hello-world
  spec:
    selector:
      matchLabels:
        app: hello-world
    replicas: 2
    template:
      metadata:
        labels:
          app: hello-world
      spec:
        containers:
          - name: hello-world
            image: gcr.io/google-samples/node-hello:1.0
            imagePullPolicy: Always
            ports:
              - containerPort: 8080
                protocol: TCP
    ---  

apiVersion: v1
kind: Service
metadata:
  name: hello-world
spec:
  selector:
    app: hello-world
  ports:
    - port: 80
      targetPort: 8080
      protocol: TCP
  type: LoadBalancer
  externalTrafficPolicy: Cluster

```



Sep 02, 2019



Aug 01, 2019



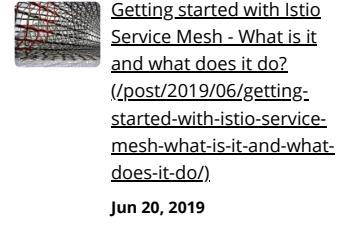
Jul 22, 2019



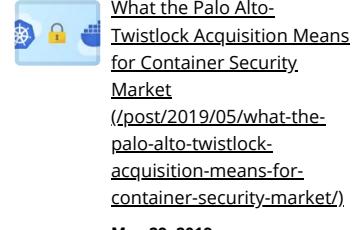
Jun 25, 2019



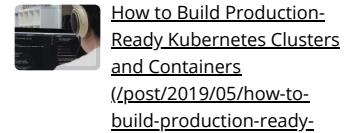
Jun 24, 2019



Jun 20, 2019

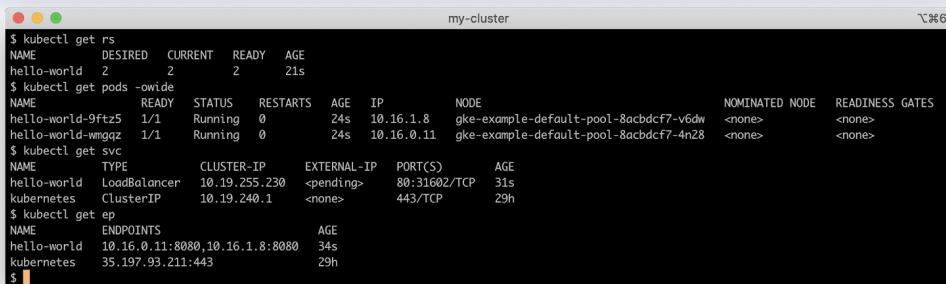


May 29, 2019



These manifests should result in the creation of two pods as part of the `hello-world` ReplicaSet, and a `hello-world` service resource with an external-facing load balancer, if the cloud provider and cluster network supports it. It should also create a Kubernetes Endpoint resource with two entries in the `host:port` notation, one for each of the pods, with the pod IP as the host value and port 8080.

(/)



```
$ kubectl get rs
NAME      DESIRED  CURRENT  READY  AGE
hello-world  2        2        2     21s

$ kubectl get pods -owide
NAME           READY  STATUS   RESTARTS  AGE  IP          NODE  NOMINATED-NODE  READINESS-GATES
hello-world-9ftz5  1/1   Running  0         24s  10.16.1.8  gke-example-default-pool-8acbdcf7-v6dw  <none>  <none>
hello-world-wmgqz 1/1   Running  0         24s  10.16.0.11  gke-example-default-pool-8acbdcf7-4n28  <none>  <none>

$ kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
hello-world  LoadBalancer  10.19.255.230  <pending>  80:31602/TCP  31s
kubernetes  ClusterIP  10.19.240.1  <none>       443/TCP     29h

$ kubectl get ep
NAME      ENDPOINTS      AGE
hello-world  10.16.0.11:8080,10.16.1.8:8080  34s
kubernetes  35.197.93.211:443  29h
```

For reference, our cluster has the following IP networks:

Node - 10.138.15.0/24

Cluster - 10.16.0.0/14

Service - 10.19.240.0/20

Our service has a Virtual IP address (VIP) of 10.19.240.1 in the cluster CIDR block.

We are now ready to follow the request's journey into the Kubernetes cluster, beginning at the load balancer.

The Load Balancer

While Kubernetes, both natively and through ingress controllers, offers a number of ways to expose a service, we will use the standard Service resource of type LoadBalancer. Our hello-world service needs a [GCP network load balancer](#). Every GKE cluster has a [cloud controller](#), which interfaces between the cluster and the API endpoints for GCP services needed to create cluster resources automatically, including our load balancer. (All cloud providers offer different classes of load balancers with varying options and characteristics.)

To see where the external load balancer fits in, first we need to look at the cluster from a different viewpoint.



[Guide to Kubernetes](#)

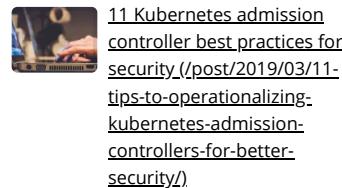
[INGRESS NETWORK POLICIES](#)
[\(/post/2019/04/setting-up-kubernetes-network-policies-a-detailed-guide/\)](#)

Apr 05, 2019



[New Kubernetes Security Vulnerabilities Disclosed: CVE-2019-1002101 and CVE-2019-9946](#)
[\(/post/2019/03/new-kubernetes-security-vulnerabilities-discovered-cve-2019-1002101-and-cve-2019-9946/\)](#)

Mar 29, 2019



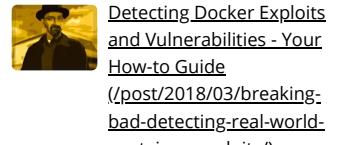
[11 Kubernetes admission controller best practices for security](#)
[\(/post/2019/03/11-tips-to-operationalizing-kubernetes-admission-controllers-for-better-security/\)](#)

Mar 22, 2019



[7 Critical Kubernetes Security Issues Resolved by Upgrading Your k8s](#)
[\(/post/2019/01/critical-kubernetes-security-issues-resolved-in-recent-kubernetes-versions/\)](#)

Jan 03, 2019



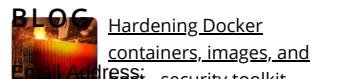
[Detecting Docker Exploits and Vulnerabilities - Your How-to Guide](#)
[\(/post/2018/03/breaking-bad-detecting-real-world-container-exploits/\)](#)

Mar 08, 2018



[Docker Forensics for Containers: How to Conduct Investigations](#)
[\(/post/2017/08/csi-container-edition-forensics-in-the-age-of-containers/\)](#)

[SUBSCRIBE TO OUR](#)



[BLOG](#) [Hardening Docker containers, images, and host - security toolkit](#)

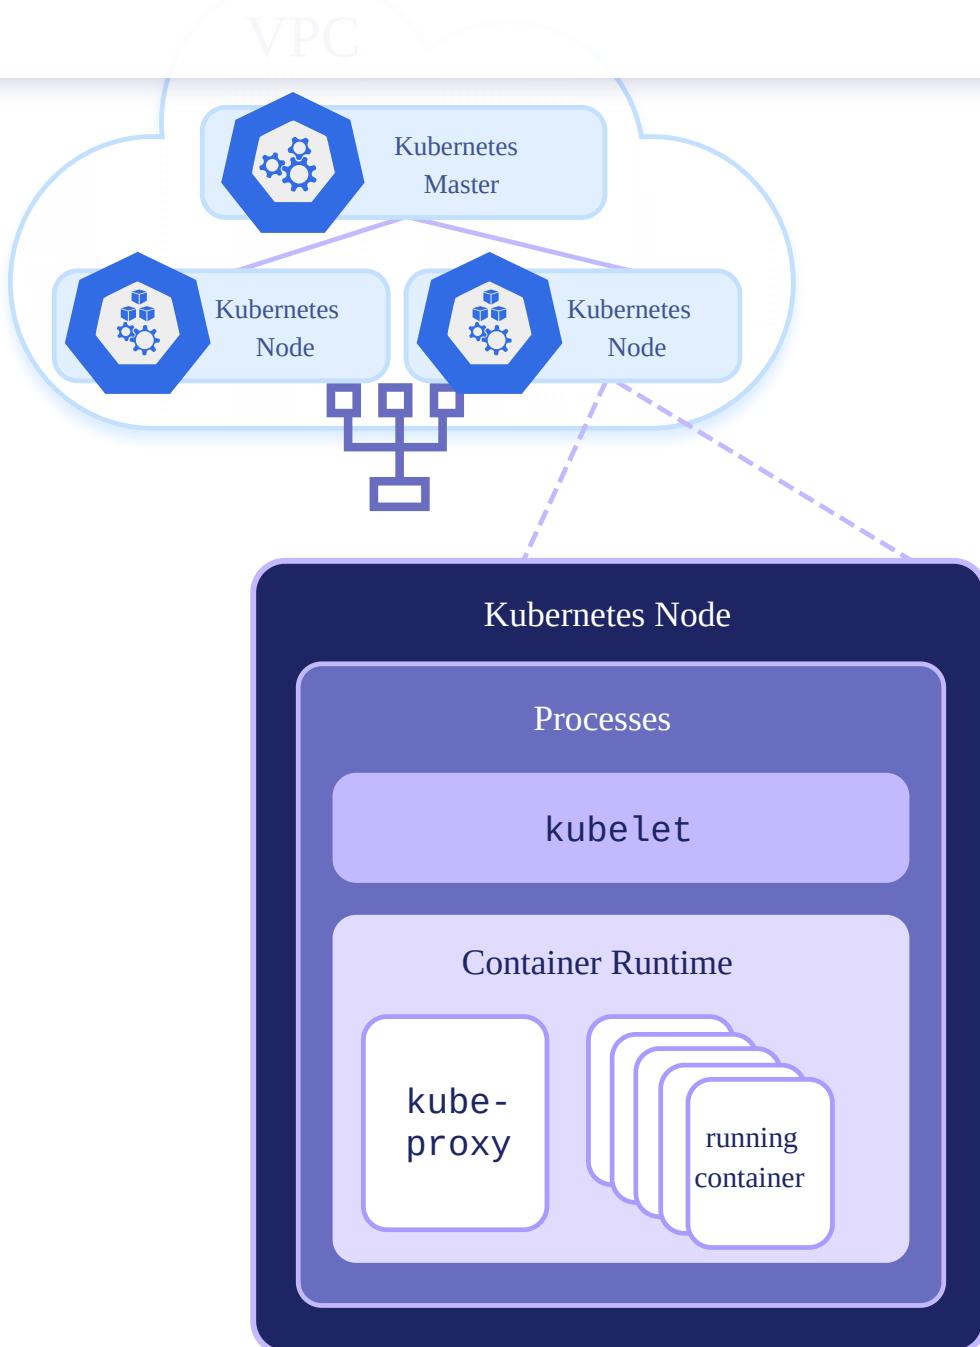
[hosts-against-](#)

[urity-](#)

[SUBSCRIBE ME](#)

Aug 10, 2017

(1)



kube-proxy

Each node has a kube-proxy container process. (In the Kubernetes frame of reference, that kube-proxy container is in a pod in the kube-system namespace.) kube-proxy manages forwarding of traffic addressed to the virtual IP addresses (VIPs) of the cluster's Kubernetes Service objects to the appropriate backend pods. kube-proxy currently supports three different operation modes:

User space: This mode gets its name because the service routing takes place in kube-proxy in the user process space instead of in the kernel network stack. It is not commonly used as it is slow and

les: This mode uses Linux kernel-level Netfilter rules to configure all routing for Kubernetes

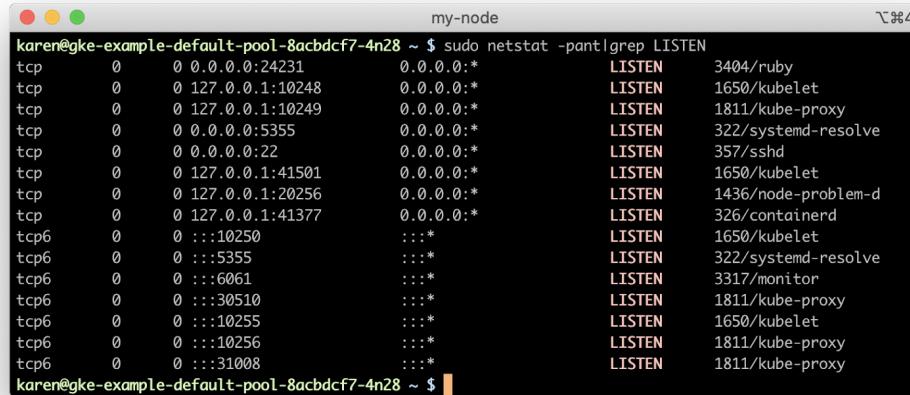
(/) Services. This mode is the default for kube-proxy on most platforms. When load balancing for

multiple backend pods, it uses unweighted round-robin scheduling.

IPVS (IP Virtual Server): Built on the Netfilter framework, IPVS implements Layer-4 load balancing in the Linux kernel, supporting multiple load-balancing algorithms, including least connections and shortest expected delay. This kube-proxy mode became generally available in Kubernetes 1.11, but it requires the Linux kernel to have the IPVS modules loaded. It is also not as widely supported by various Kubernetes networking projects as the iptables mode.

kube-proxy in our GKE cluster runs in iptables mode, so we will look at how that mode works.

If we look at the `hello-world` service we created, we can see that it has been assigned a node port (a network port for the node's IP address) of 30510. Dynamically-assigned ports on the node network allow multiple Kubernetes services hosted in the cluster to use the same Internet-facing port in their endpoints. If our service had been deployed to a standard Amazon Elastic Kubernetes Service (EKS) cluster, it would be served by an Elastic Load Balancer which would send incoming connections to our service's node port on nodes with a live service pod. However, Google Cloud Platform (GCP) network load balancers only forward traffic to the targets on the same port as the incoming port on the load balancer, i.e., traffic to port 80 on the load balancer will be sent to port 80 on the target backend instance. The `hello-world` pods are definitely not listening on port 80 of the node. If we run `netstat` on the node, we see that no process is listening on that port.



```

my-node
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $ sudo netstat -antl | grep LISTEN
tcp        0      0 0.0.0.0:24231          0.0.0.0:*              LISTEN      3404/ruby
tcp        0      0 127.0.0.1:10248        0.0.0.0:*              LISTEN      1650/kubelet
tcp        0      0 127.0.0.1:10249        0.0.0.0:*              LISTEN      1811/kube-proxy
tcp        0      0 0.0.0.0:5355           0.0.0.0:*              LISTEN      322/systemd-resolve
tcp        0      0 0.0.0.0:22            0.0.0.0:*              LISTEN      357/sshd
tcp        0      0 127.0.0.1:41501        0.0.0.0:*              LISTEN      1650/kubelet
tcp        0      0 127.0.0.1:20256        0.0.0.0:*              LISTEN      1436/node-problem-d
tcp        0      0 127.0.0.1:41377        0.0.0.0:*              LISTEN      326/containerd
tcp6       0      0 ::1:10250            ::*:                  LISTEN      1650/kubelet
tcp6       0      0 ::1:5355             ::*:                  LISTEN      322/systemd-resolve
tcp6       0      0 ::1:6061             ::*:                  LISTEN      3317/monitor
tcp6       0      0 ::1:30510            ::*:                  LISTEN      1811/kube-proxy
tcp6       0      0 ::1:10255            ::*:                  LISTEN      1650/kubelet
tcp6       0      0 ::1:10256            ::*:                  LISTEN      1811/kube-proxy
tcp6       0      0 ::1:31008            ::*:                  LISTEN      1811/kube-proxy
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $

```

So how do requests through the load balancer make a successful connection? If kube-proxy were running in the user space mode, it would actually be proxying connections to backend pods. In iptables mode, though, kube-proxy configures Netfilter chains so the connection is routed directly to the backend container's endpoint by the node's kernel.

(1)

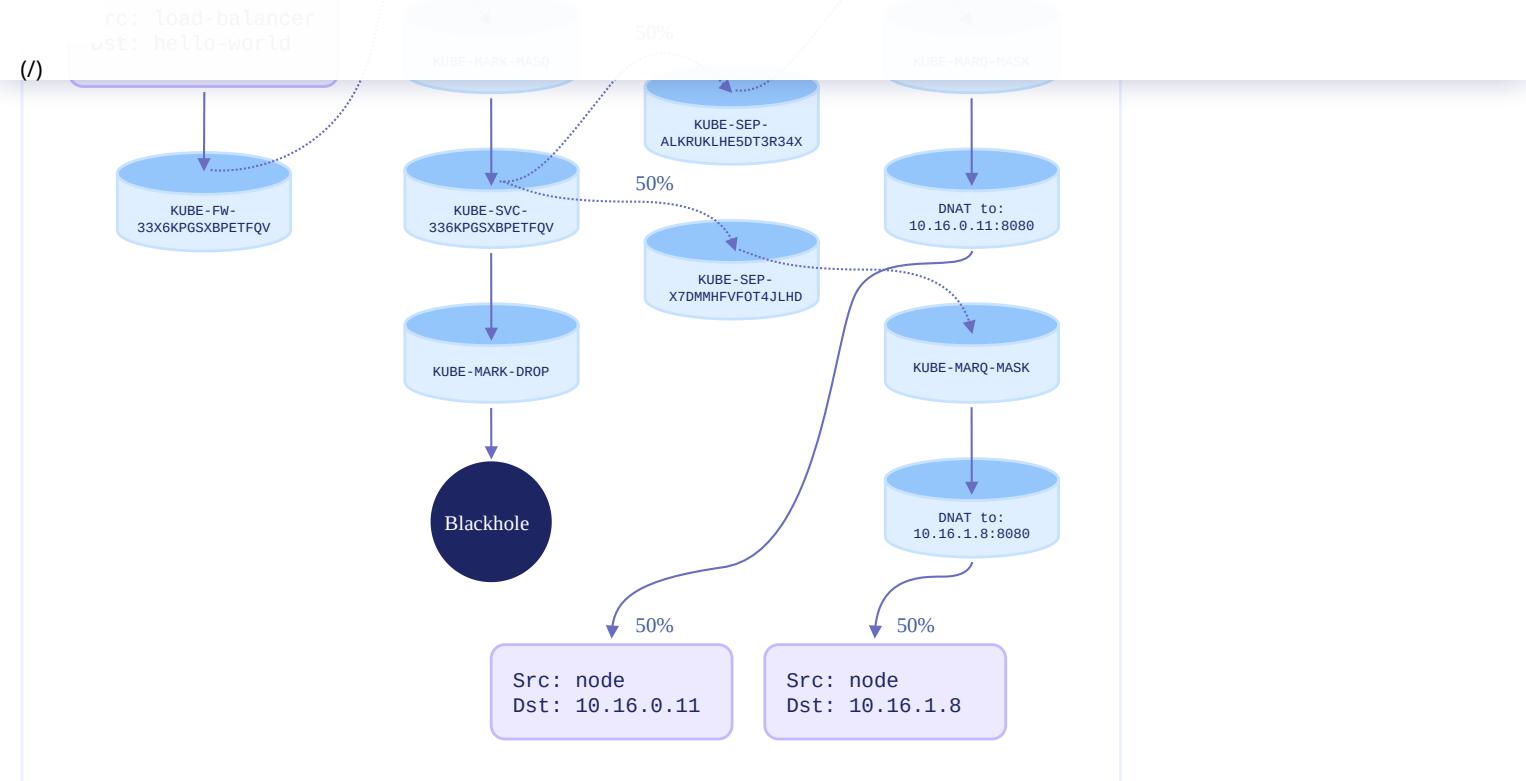
In GKE cluster, if we log in to one of the nodes and run iptables, we can see these rules.

```
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $ sudo iptables -t nat -L KUBE-SERVICES
Chain KUBE-SERVICES (2 references)
target  prot opt source          destination
KUBE-MARK-MASQ  tcp  --  !10.16.0.0/14  10.19.248.159    /* kube-system/metrics-server: cluster IP */ & tcp dpt:https
KUBE-SVC-LC5QY6WV2JU6NZ  tcp  --  anywhere    10.19.248.159    /* kube-system/metrics-server: cluster IP */ & tcp dpt:https
KUBE-MARK-MASQ  tcp  --  !10.16.0.0/14  10.19.255.230   /* default/hello-world: cluster IP */ & tcp dpt:http
KUBE-SVC-33X6KPGSXBPETFQV  tcp  --  anywhere    10.19.255.230   /* default/hello-world: cluster IP */ & tcp dpt:http
KUBE-FW-33X6KPGSXBPETFQV  tcp  --  anywhere    157.119.83.34.bc.googleusercontent.com /* default/hello-world: loadbalancer IP */ & tcp dpt:http
KUBE-MARK-MASQ  tcp  --  !10.16.0.0/14  10.19.246.1    /* default/kubernetes:https cluster IP */ & tcp dpt:https
KUBE-SVC-4M4PHTMRTRNKG  tcp  --  anywhere    10.19.246.1    /* default/kubernetes:https cluster IP */ & tcp dpt:https
KUBE-MARK-MASQ  udp  --  !10.16.0.0/14  10.19.241.249   /* kube-system:cluster IP */ & udp dpt:http
KUBE-SVC-XPM0J6VSLQALMWS  tcp  --  anywhere    10.19.241.249   /* kube-system/http-backend:tcp cluster IP */ & tcp dpt:http
KUBE-MARK-MASQ  udp  --  !10.16.0.0/14  10.19.240.10   /* kube-system/kube-dns:dns cluster IP */ & udp dpt:domain
KUBE-MARK-MASQ  tcp  --  !10.16.0.0/14  10.19.240.10   /* kube-system/kube-dns:dns-tcp cluster IP */ & tcp dpt:domain
KUBE-SVC-TCOU7-COKEZGVNU  udp  --  anywhere    10.19.240.10   /* kube-system/default:udp-backend:tcp cluster IP */ & udp dpt:domain
KUBE-MARK-MASQ  tcp  --  !10.16.0.0/14  10.19.240.10   /* kube-system/kube-dns:cluster IP */ & tcp dpt:domain
KUBE-SVC-ERIFXISQEPP7F04  tcp  --  anywhere    10.19.240.10   /* kube-system/kube-dns-tcp:cluster IP */ & tcp dpt:domain
KUBE-MARK-MASQ  tcp  --  !10.16.0.0/14  10.19.246.96   /* kube-system/heapster: cluster IP */ & tcp dpt:http
KUBE-SVC-BJMGW63USR2HCRZ  tcp  --  anywhere    10.19.246.96   /* kube-system/heapster: cluster IP */ & tcp dpt:http
KUBE-NODEPORTS  all  --  anywhere    anywhere      /* kubernetes service:nodeports; NOTE: this must be the last rule in this chain */ ADDRTYPE match dst-type LOCAL
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $
```

Thanks to the rule comments, we can get the name of the filter chain that matches incoming connections from the service's load balancer to our `hello-world` service and follow that chain's rules. (In the absence of a rule comment, we still could have matched the rule's source IP address to the service's load balancer.)

```
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $ sudo iptables -t nat -L KUBE-FW-33X6KPGSXBPETFQV
Chain KUBE-FW-33X6KPGSXBPETFQV (1 references)
target  prot opt source          destination
KUBE-MARK-MASQ  all  --  anywhere    anywhere      /* default/hello-world: loadbalancer IP */
KUBE-SVC-33X6KPGSXBPETFQV  all  --  anywhere    anywhere      /* default/hello-world: loadbalancer IP */
KUBE-MARK-DROP  all  --  anywhere    anywhere      /* default/hello-world: loadbalancer IP */
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $ sudo iptables -t nat -L KUBE-MARK-MASQ
Chain KUBE-MARK-MASQ (20 references)
target  prot opt source          destination
MARK  all  --  anywhere    anywhere      MARK or 0x4000
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $ sudo iptables -t nat -L KUBE-SVC-33X6KPGSXBPETFQV
Chain KUBE-SVC-33X6KPGSXBPETFQV (3 references)
target  prot opt source          destination
KUBE-SEP-ALRUKLHE5DT3R34X  all  --  anywhere    anywhere      statistic mode random probability 0.500000000000
KUBE-SEP-X7DMMHFVFOT4JLHD  all  --  anywhere    anywhere
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $ sudo iptables -t nat -L KUBE-MARK-DROP
Chain KUBE-MARK-DROP (1 references)
target  prot opt source          destination
MARK  all  --  anywhere    anywhere      MARK or 0x8000
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $ sudo iptables -t nat -L KUBE-SEP-ALRUKLHE5DT3R34X
Chain KUBE-SEP-ALRUKLHE5DT3R34X (1 references)
target  prot opt source          destination
KUBE-MARK-MASQ  all  --  10.16.0.11   anywhere
DNAT  tcp  --  anywhere    anywhere      tcp to:10.16.0.11:8080
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $ sudo iptables -t nat -L KUBE-SEP-X7DMMHFVFOT4JLHD
Chain KUBE-SEP-X7DMMHFVFOT4JLHD (1 references)
target  prot opt source          destination
KUBE-MARK-MASQ  all  --  10.16.1.8    anywhere
DNAT  tcp  --  anywhere    anywhere      tcp to:10.16.1.8:8080
karen@gke-example-default-pool-8acbdcf7-4n28 ~ $
```

We can also visualize the chains and rules used in the network stack for evaluating and modifying the packet to see how the service we created in our cluster directs traffic to the replica set members.



The KUBE-FW-33X6KPGSXBPFQV chain has three rules, each adding another chain for handling the packet.

1. KUBE-MARK-MASQ adds a Netfilter mark to packets destined for the `hello-world` service which originate outside the cluster's network. Packets with this mark will be altered in a POSTROUTING rule to use source network address translation (SNAT) with the node's IP address as their source IP address.
2. The KUBE-SVC-33X6KPGSXBPFQV chain applies to all traffic bound for our `hello-world` service, regardless of source, and has rules for each of the service endpoints (the two pods, in this case). Which endpoint chain to use gets determined in a purely random fashion.
 1. KUBE-MARK-MASQ again adds a Netfilter mark to the packet for SNAT, if needed
 2. The DNAT rule sets up a destination NAT using the `10.16.0.11:8080` endpoint as the destination.
2. KUBE-SEP-X7DMMHFVF0T4JLHD
 1. KUBE-MARK-MASQ again adds a Netfilter mark to the packet for SNAT, if needed
 2. The DNAT rule sets up a destination NAT using the `10.16.1.8:8080` endpoint as the destination.
3. KUBE-MARK-DROP adds a Netfilter mark to packets which do not have destination NAT enabled by this point. These packets will be discarded in the KUBE-FIREWALL chain.

to preference for routing to the pod on the node that receives the request from the Cloud Load

balancer. If we change the `externalTrafficPolicy` in the service spec to `Local`, however, that

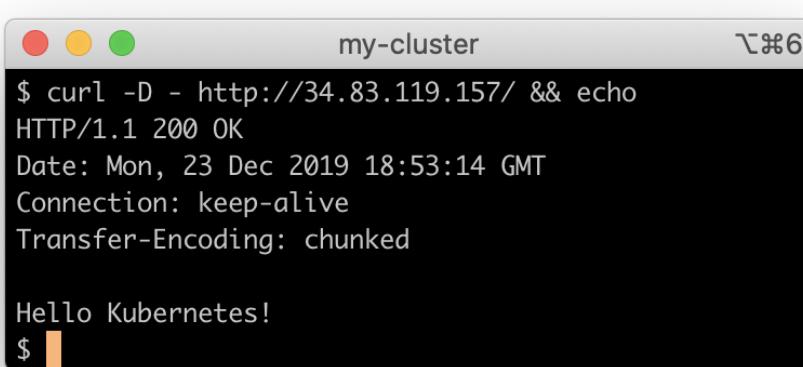
(/)
would change. Not only would the request go to a pod on the node receiving the request, if one exists, but it also means a node without a service pod will refuse the connection. Therefore the `Local` policy generally needs to be used with Kubernetes daemon sets, which schedule a pod on each node in the cluster. While specifying local delivery would obviously reduce the mean network latency for requests, it can lead to uneven load across the service's pods.

The Pod Network

This post will not dive into details on pod networking, but in our GKE cluster, the pod network has its own CIDR block separate from the node's network. [The Kubernetes network model](#) requires all pods in the cluster to be able to address each other directly, regardless of their host node. GKE clusters use the kubenet CNI, which creates network bridge interfaces to the pod network on each node, giving each node its own dedicated CIDR block of pod IP addresses to simplify allocation and routing. The Google Compute Engine (GCE) network can route this pod network traffic between VMs.

The Request

And this is how we get our HTTP 200 response code.



```
$ curl -D - http://34.83.119.157/ && echo
HTTP/1.1 200 OK
Date: Mon, 23 Dec 2019 18:53:14 GMT
Connection: keep-alive
Transfer-Encoding: chunked

Hello Kubernetes!
$
```

Routing Variables

This post mentioned some of the ways different options offered by various Kubernetes platforms can change routing. Here is a non-comprehensive list:

Example. Examples in Amazon EKS would have looked very different, because the AWS VPC CNI

- (/)
- places pods directly on the nodes' VPC network.

Kubernetes Network Policy: One of the most popular CNI plugins implementing network policies, Calico, creates a virtual network interface on the nodes for each pod and uses Netfilter rules to enforce its firewall rules.

The kube-proxy IPVS routing mode moves the service routing and NATing out of the Netfilter rules, for the most part, although it does still make use of Netfilter.

External load balancers or other sources which can send traffic directly to the service's node port would match a different chain (KUBE-NODEPORTS) in iptables.

Kubernetes [ingress controllers](#) may change edge service routing in a number of ways.

Service meshes like Istio may bypass kube-proxy and make direct connections for internal routing between service pods.

Securing Services

No universal method for adding firewall restrictions to cloud load balancers that are created by Kubernetes Service resources exists. Some cloud providers honor the [loadBalancerSourceRanges](#) field in the Service spec, which allows you to provide a whitelist of IP CIDR blocks allowed to connect to the load balancer. If a cloud provider does not honor this field, it will be silently ignored, so take care to verify the network configuration of your external load balancers. For providers that do not support the `LoadBalancerSourceRanges` field, you should assume your service endpoints on the load balancers will be open to the world unless you take action at the cloud provider level to lock down the load balancers and the cloud networks on which they run. The default firewall settings for cloud provider load balancer offerings vary wildly and depend on many factors. Some cloud providers may also support annotations to the Service object to configure load balancer security.

Note that we did not install the Calico CNI by enabling Kubernetes Network Policy support in our GKE cluster, because Calico creates a large number of additional iptables rules, adding extra steps when visually tracing virtual routing to a pod. However, we strongly recommend using a CNI that implements the NetworkPolicy API in production clusters and creating [policies that restrict pod traffic](#) (<https://www.stackrox.com/post/2019/04/setting-up-kubernetes-network-policies-a-detailed-guide/>).

Pods created with the `HostNetwork` attribute enabled will share the node's network space. While some valid use cases exist for doing so, generally most pods do not need to be on the host network, and particularly for pods running with root privileges, it could allow a compromised container to sniff network traffic. If you need to expose a container port on the node's network and using a

For a full overview of the Container Security Maturity Model, see the Container Security Maturity Model - Understand your security needs at every stage of the container journey.

(I) Pods using the host network should not run with the NET_ADMIN capability, which would allow them

to read and modify the node's firewall rules.

The Kubernetes network requires a large number of moving pieces. It is quite complicated, but having a basic understanding of what is taking place in your cluster will help you more effectively monitor, secure, and protect it.

Sources and Further Reading

<https://kubernetes.io/docs/concepts/services-networking/service/>

<https://kubernetes.io/docs/concepts/cluster-administration/networking/>

<https://twitter.com/thockin/status/1191766983735296000>

<https://kubernetes.io/blog/2018/07/09/ipv6-based-in-cluster-load-balancing-deep-dive/>

<https://netfilter.org/documentation/HOWTO/NAT-HOWTO-6.html>

Categories:

[Kubernetes](#) ([/Categories/Kubernetes](#))

Tags:

[Kubernetes Networking](#) ([/Tags/Kubernetes-Networking](#))

100 View Street, Suite 204
Mountain View, CA 94041
+1 (650) 385-8329

WHY	USE CASES	ENVIRONMENTS	FEATURED RESOURCES
STACKROX (/WHY- STACKROX/)	Visibility (/use- cases/visibility/) Vulnerability	AWS (/solutions/aws- security/) Azure	Kubernetes Security 101 (https://www.stackrox.com/post/2019/07/kubernetes-security-101/)
PLATFORM (/PLATFORM/)	Management (/use- cases/vulnerability- management/)	(/solutions/microsoft- azure-security/)	Docker Security 101 (https://www.stackrox.com/post/2019/09/docker-security-101/)

[CONTACT US](#)

CUSTOMERS (/CUSTOMERS/)	Compliance (/use- cases/compliance/)	Docker (/solutions/docker- security/)	EKS Security Best Practices (https://www.stackrox.com/post/2019/09/amazon-eks-security-best-practices/)
RESOURCES (/ASSETS/)	Network Segmentation (/use- cases/network- segmentation/)	Google Cloud Platform Security for google- cloud-platform/	Kubernetes Network Policies (https://www.stackrox.com/post/2019/04/setting-up-kubernetes-network-policies-a-detailed-guide/)
BLOG (/POST/)	Risk Profiling (/use- cases/risk-profiling/)		Kubernetes Configuration Best Practices (https://www.stackrox.com/post/2019/09/12- ,

(/)

Management (/use-cases/configuration-management/)	Pivotal Container Service (PKS) (/solutions/pks-security/)	Intro to Istio Service Mesh (https://www.stackrox.com/post/2019/06/getting-started-with-istio-service-mesh-what-is-it-and-what-does-it-do/)
Threat Detection (/use-cases/threat-detection/)	Istio Security Basics (https://www.stackrox.com/post/2019/08/istio-security-basics-running-microservices-on-zero-trust-networks/)	
Incident Response (/use-cases/incident-response/)	Federal Agencies (/solutions/federal-agencies/)	Kubernetes Adoption and Security Trends (https://www.stackrox.com/kubernetes-adoption-and-security-trends-and-market-share-for-containers/)
		Admission Controllers (https://www.stackrox.com/post/2019/03/11-tips-to-operationalizing-kubernetes-admission-controllers-for-better-security/)
		Docker Container Forensics (https://www.stackrox.com/post/2017/08/csi-container-edition-forensics-in-the-age-of-containers/)
		Detecting Docker Exploits and Vulnerabilities (https://www.stackrox.com/post/2017/08/csi-container-edition-forensics-in-the-age-of-containers/)
		Hardening Docker Containers and Images (https://www.stackrox.com/post/2017/08/hardening-docker-containers-and-hosts-against-vulnerabilities-a-security-toolkit/)